

Robustness Guarantees for Bayesian Inference with Gaussian Processes

Andrea Patane

Department of Computer Science, University of Oxford

Joint work with:

Luca Cardelli, Marta Kwiatkowska and Luca Laurenti

Outline

- Motivations.
- Background: Bayesian Inference with Gaussian Processes.
- Problem Formulation: Probabilistic invariance.
- Methods:
 - Safe-approximation of invariance property.
 - Branch and Bound optimisation scheme for GPs.
- Case of Study: Empirical analysis of ReLU fully-connected Neural Networks via GP with ReLU kernel.

Robustness for Bayesian Learning, Why?

- Bayesian methods are employed in **safety critical** applications, where **uncertainty** estimation is necessary (e.g. diagnosis, medicine intake, control systems...).
- Robustness guarantees are needed to prove the correctness of the model in a **probabilistic** fashion.
- Current methods either neglect **uncertainty** or are based on **empirical** approaches (e.g. variance thresholding)

Problem: Provide probabilistic guarantees for GPs.


Background

Bayesian Inference with GPs

(in Figures)

- Step 1: Definition of a GP prior distribution.


$$f \sim \mathcal{GP}$$



Bayesian Inference with GPs (in Figures)


- Step 1: Definition of a GP prior distribution.

$$f \sim \mathcal{GP}$$



- Step 2: Conditioning on training data.

$$f \sim \mathcal{GP} \mid \mathbf{x}$$



Bayesian Inference with GPs (in Formulas)

- Let z be a GP with prior mean μ and variance Σ . Consider a training set $D = \{(x_i, y_i)\}_{i=1, \dots, N}$. The goal of Bayesian inference is to find:

$$\hat{z} = z \mid D$$

- For GPs this can be done [analytically](#), obtaining a GP with posteriori mean and variance given by:

$$\hat{\mu}(x^*) = \mu(x^*) + \Sigma_{x^*, \mathcal{D}} \Sigma_{\mathcal{D}, \mathcal{D}}^{-1} (\mathbf{y} - \mu_{\mathcal{D}})$$

$$\hat{\Sigma}_{x^*, x^*} = \Sigma_{x^*, x^*} - \Sigma_{x^*, \mathcal{D}} \Sigma_{\mathcal{D}, \mathcal{D}}^{-1} \Sigma_{x^*, \mathcal{D}}^T$$

Problem Formulation

Probabilistic Invariance

- Probabilistic generalisation of problem associated with existence of [local adversarial examples](#).
- Intuitively, we want to count the number of functions extracted from the GP for which deterministic invariance does not hold.


Probabilistic Invariance

- Probabilistic generalisation of problem associated with existence of **local adversarial examples**.
- Intuitively, we want to count the number of functions extracted from the GP for which deterministic invariance does not hold.


Consider x^* and a neighbourhood T . Let δ be the adversarial threshold, then **invariance probability** is defined by:

$$\phi(x^*, T, \delta) = P(\exists x' \in T \text{ s.t. } ||\hat{z}(x') - \hat{z}(x^*)|| > \delta)$$



Probabilistic Invariance (in Figures)



Probabilistic Invariance (in Figures)





$f_1 \sim \mathcal{GP}$




Safe

$f_2 \sim \mathcal{GP}$





Unsafe

Probabilistic Invariance (in Figures)





$f_1 \sim \mathcal{GP}$



Safe

$f_2 \sim \mathcal{GP}$




Unsafe

•
•
•



•
•
•

•
•
•

Probabilistic Invariance (in Figures)





$f_1 \sim \mathcal{GP}$



Safe

$f_2 \sim \mathcal{GP}$



Unsafe

•

•

•

•

•



•

•

•


•

$f_\infty \sim \mathcal{GP}$





Safe

Probabilistic Invariance (in Figures)





$f_1 \sim \mathcal{GP}$



Safe

$f_2 \sim \mathcal{GP}$



Unsafe

•

•

•



•
•
•

•

•

•

$f_\infty \sim \mathcal{GP}$



Safe

$\phi(x^*, T, \delta)$

Methods

Upper-bound on Invariance

- Computation of $\phi(x^*, T, \delta)$ is far from trivial as it involves solutions of **uncountably** many optimisation problems. Instead, we compute a **safe approximation**.

Upper-bound on Invariance

- Computation of $\phi(x^*, T, \delta)$ is far from trivial as it involves solutions of **uncountably** many optimisation problems. Instead, we compute a **safe approximation**.

Theorem 1: For every output dimension i let:

$$\eta_i = \frac{\delta - \sup_{x \in T} |\mu^o(x^*, x)|_1}{n} - 12 \int_0^{\frac{1}{2} \sup_{x_1, x_2 \in T} d_{x^*}^{(i)}(x_1, x_2)} \sqrt{\ln \left(\left(\frac{\sqrt{m} K_{x^*}^{(i)} D}{z} + 1 \right)^m \right)} dz$$

Then:

$$\phi(x^*, T, \delta | \mathcal{D}) \leq \hat{\phi}(x^*, T, \delta | \mathcal{D}) := 2 \sum_{i=1}^n e^{-\frac{\bar{\eta}_i^2}{2\xi^{(i)}}}$$

Upper-bound on Invariance

Theorem 1: For every output dimension i let:

Maximum mean
difference

$$\eta_i = \frac{\delta - \sup_{x \in T} |\mu^o(x^*, x)|_1}{n} - 12 \int_0^{\frac{1}{2} \sup_{x_1, x_2 \in T} d_{x^*}^{(i)}(x_1, x_2)} \sqrt{\ln \left(\left(\frac{\sqrt{m} K_{x^*}^{(i)} D}{z} + 1 \right)^m \right)} dz$$

Then:

$$\phi(x^*, T, \delta | \mathcal{D}) \leq \hat{\phi}(x^*, T, \delta | \mathcal{D}) := 2 \sum_{i=1}^n e^{-\frac{\bar{\eta}_i^2}{2\xi^{(i)}}}$$

Upper-bound on Invariance

Theorem 1: For every output dimension i let:

$$\eta_i = \frac{\delta - \overbrace{\sup_{x \in T} |\mu^o(x^*, x)|_1}^{\text{Maximum mean difference}}}{n \underbrace{\text{Lipschitz constant on Variance}}} - 12 \int_0^{\frac{1}{2} \sup_{x_1, x_2 \in T} d_{x^*}^{(i)}(x_1, x_2)} \sqrt{\ln \left(\left(\frac{\sqrt{m} \overbrace{K_{x^*}^{(i)}}^{\text{Lipschitz constant on Variance}} D}{z} + 1 \right)^m \right)} dz$$

Then:

$$\phi(x^*, T, \delta | \mathcal{D}) \leq \hat{\phi}(x^*, T, \delta | \mathcal{D}) := 2 \sum_{i=1}^n e^{-\frac{\bar{\eta}_i^2}{2\xi^{(i)}}}$$

Upper-bound on Invariance

Theorem 1: For every output dimension i let:

$$\eta_i = \frac{\delta - \overbrace{\sup_{x \in T} |\mu^o(x^*, x)|_1}^{\text{Maximum mean difference}}}{n \underbrace{\text{Lipschitz constant on Variance}}} - 12 \int_0^{\frac{1}{2} \sup_{x_1, x_2 \in T} d_{x^*}^{(i)}(x_1, x_2)} \sqrt{\ln \left(\left(\frac{\sqrt{m} \overbrace{K_{x^*}^{(i)}}^{\text{Maximum correlation}} D}{z} + 1 \right)^m \right)} dz$$

Then:

$$\phi(x^*, T, \delta | \mathcal{D}) \leq \hat{\phi}(x^*, T, \delta | \mathcal{D}) := 2 \sum_{i=1}^n e^{-\frac{\bar{\eta}_i^2}{2 \underbrace{\xi^{(i)}}_{\text{Maximum correlation}}}}$$

Proof Sketch

- We want to upper-bound:

$$\phi(x^*, T, \delta | \mathcal{D}) = P(\sup_{x \in T} \|z(x) - z(x^*)\| > \delta)$$

- Since $z^o(x^*, x) = z(x^*) - z(x)$ is still a GP we can employ the Borell-TIS inequality, which upper-bounds the supremum:

$$P(\sup_{x \in T} \|z^o(x^*, x)\| > \delta) \leq e^{\frac{(\delta - E[\sup_{x \in T} z^o(x^*, x)])^2}{2\sigma_T^2}}$$

- Finally, $E[\sup_{x \in T} z^o(x^*, x)]$ can be over-approximated using the Dudley entropy integral.

Constant Computation

- The **upper-bound** computation requires computation of different constants e.g.:

$$\sup_{x \in T} \mu(x^*) - \mu(x) = \mu(x^*) - \inf_{x \in T} \mu(x) = \mu(x^*) - \inf_{x \in T} \Sigma_{x, \mathcal{D}} \Sigma_{\mathcal{D}, \mathcal{D}}^{-1} \mathbf{y}$$

- We define two functions φ and ψ that **decompose** the GP variance as: $\Sigma_{x, x_i} = \psi(\varphi(x, x_i))$.
- Using **interval analysis** on φ and **optimising** ψ we can compute lower and upper bounds on each Σ_{x, x_i}
- Thanks to linearity, we propagate these to get bounds on the sup; and refine via **Branch and Bound**.

Case of Study







GPs and Neural Networks: Experimental Settings

- Bayesian fully-connected neural networks converge in distribution to specific GPs, as the number of **neurons** approaches **infinity***.
- We can employ the method we developed to perform empirical analysis of **fully connected NNs**.
- We focus on **ReLU** NNs applied to the MNIST dataset.
- For scalability, we provide **feature-level** analysis using SIFT.

*Neal, R. M. Bayesian learning for neural networks. Springer, 2012.


Parametric Analysis on Adversarial Thresholds

Feature 1 Feature 2 Feature 3 Feature 4 Feature 5



Parametric Analysis on Variance

Analysis of how variance changes in T depending on number of training samples and layers.



Conclusions

- We developed a formal approach for **invariance analysis** of Bayesian inference with Gaussian Processes.
- Developed an algorithmic approach for computation of upper-bound on invariance probability.
- We relied on the relationship between **Bayesian NNs and GPs**, to analyse NN behaviour at infinity width limit.
- Provided experimental results on MNIST.